

FindAll: A Local Search Engine for Mobile Phones

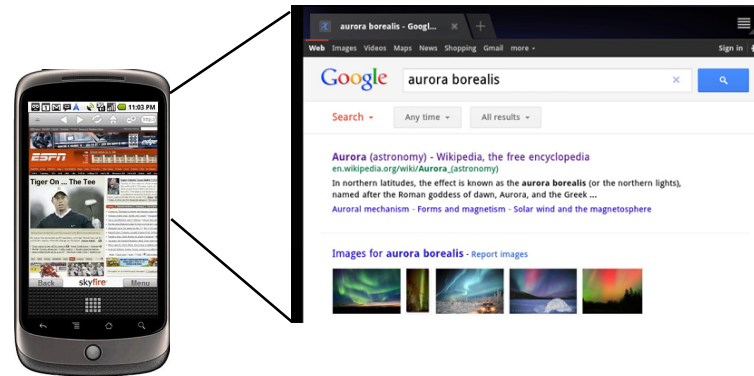
Aruna Balasubramanian, Niranjana Balasubramanian,
Sam Huston, Donald Metzger, David Wetherall



UMassAmherst



Motivation: Mobile web search depends heavily on connectivity




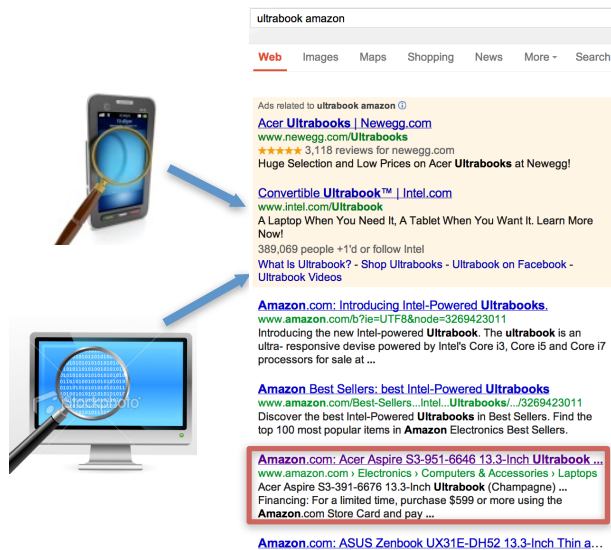
- But 3G connectivity is often poor (and expensive), or may not be available
 - Search is $>10X$ slower than on desktop

Idea: Improve web search with better re-finding

Re-finding is very common

1) Visit a page on mobile (or desktop)

2) Later, search for same page on mobile

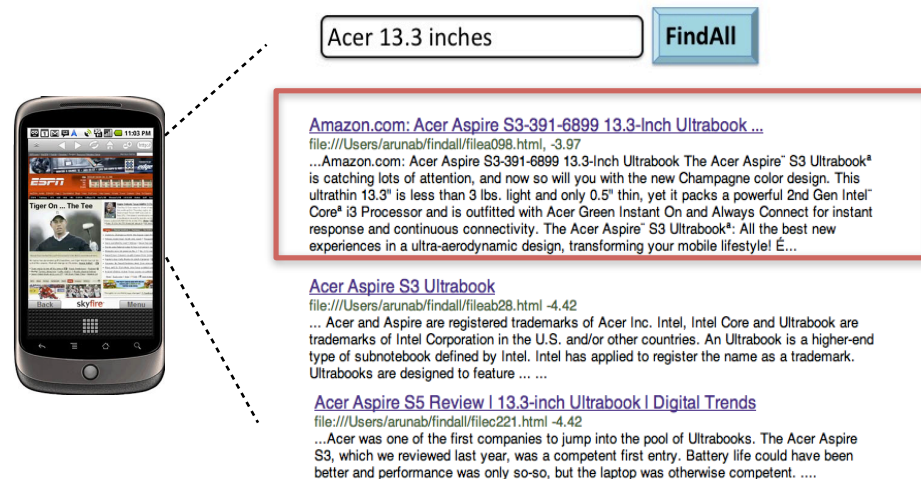


[Amazon.com: Acer Aspire S3-951-6646 13.3-Inch Ultrabook ...](#)
www.amazon.com › Electronics › Computers & Accessories › Laptops
Acer Aspire S3-951-6646 13.3-Inch **Ultrabook** (Champagne) ...
Financing: For a limited time, purchase \$599 or more using the **Amazon.com** Store Card and pay ...

- Mobile: 70% of searches for 50% of users
- Non-Mobile: 40% to 60% of all searches

Our Solution: FindAll

- Local web search on mobile for re-finding with a Google/Bing-like search interface



Trades compute/storage for connectivity

- Will show lower search latency and improved availability with reduced energy cost

Why a local search engine?

- Search with indexes (ranking) is very effective compared to alternatives:
 - Browsing history is cumbersome
 - Keyword matching on cached pages is cheap and fast, but not very effective
 - Database of search queries misses query changes and non-searched pages

Challenge: search is memory/energy intensive

Talk & Contributions

1. User study

- Identifies re-finding behaviors

2. FindAll

- Design of search engine for mobile re-finding

3. Evaluation

- Results of the tradeoffs in practice

User study

- Monitored 23 participants for 1 month
 - Grad and under-grad students
- Collected logs from user's mobile/desktop
 - Visited URL and search query (anonymized)
- Mark URL re-found if:
 - Page revisited via search query, and unchanged

Examples

Re-finding

→ URL: <http://conferences.sigcomm.org/co-next/2012/>
...
→ Search query: "conext 2012"
URL: <http://conferences.sigcomm.org/co-next/2012/>
...
→ Search query: "networking conference nice france"
URL: <http://conferences.sigcomm.org/co-next/2012/>



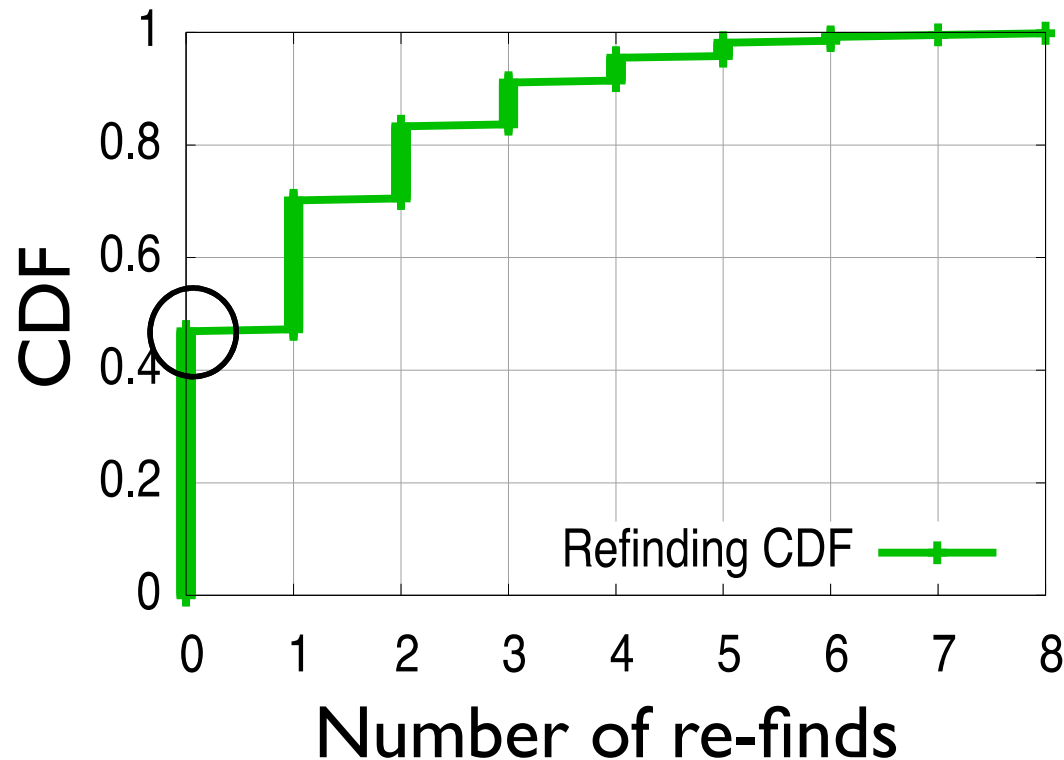
Not Re-finding



Search query: "weather"
URL: www.weather.com
...
Search query: "weather"
URL: www.weather.com
...
URL: http://wikipedia.org/wiki/J._K._Rowling
...
URL: http://wikipedia.org/wiki/J._K._Rowling

(Conservative rules, under-estimate re-finding)

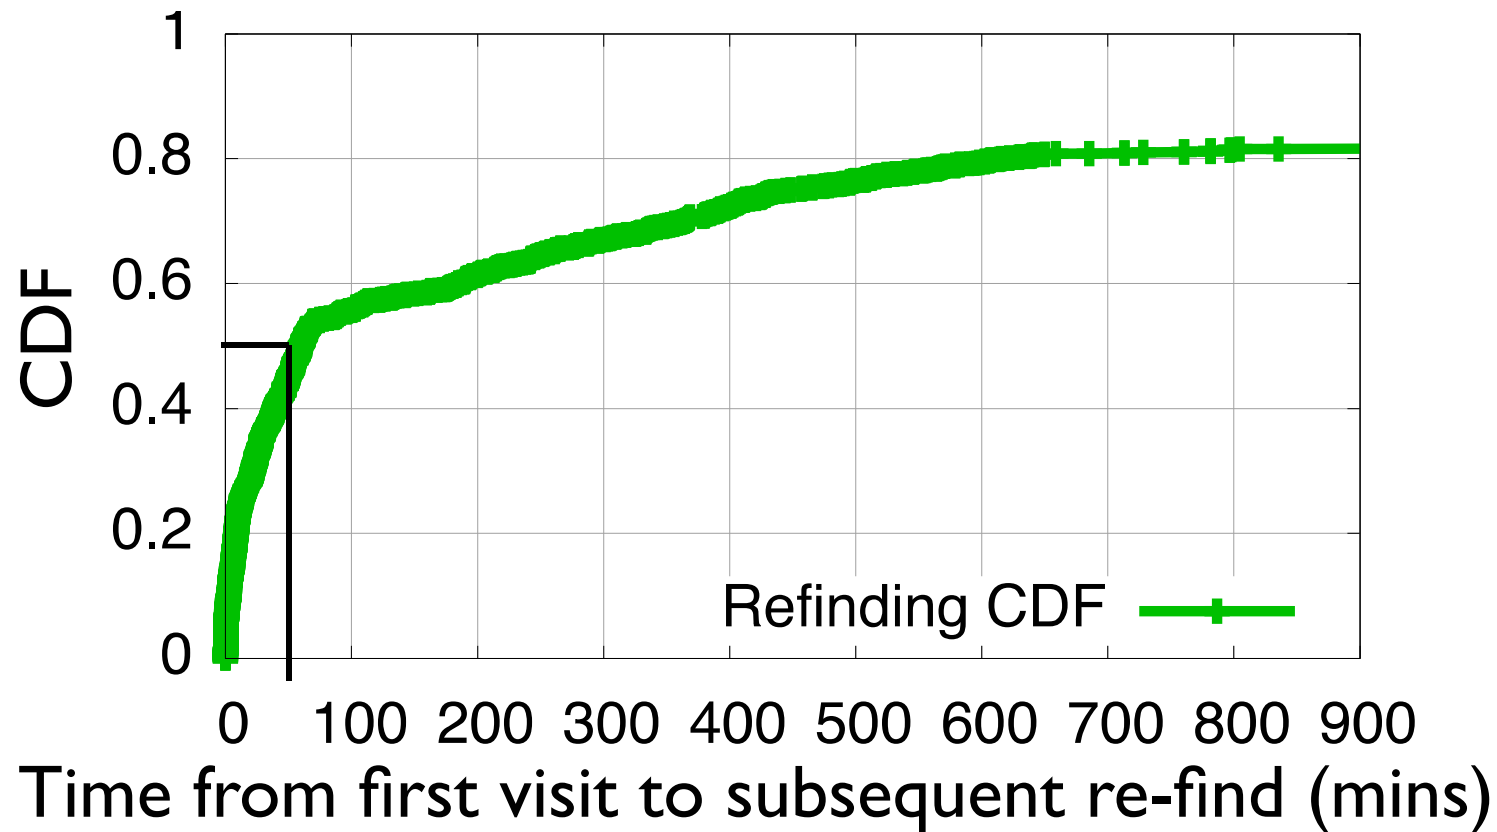
Most search is re-finding



> 20% of
re-finds
have
different
query

Lots of opportunity for improvement!


Most re-finds happen quickly



Need to index near when page is first accessed


Diverse behavior across users

Varying numbers of pages, fraction of re-finding



Re-finding behavior is fairly consistent for a given user

Sample user with avg. 43% re-finding with 9% std dev.



Can adapt to user behavior!

Outline

1. User study

- Identifies key re-finding behaviors

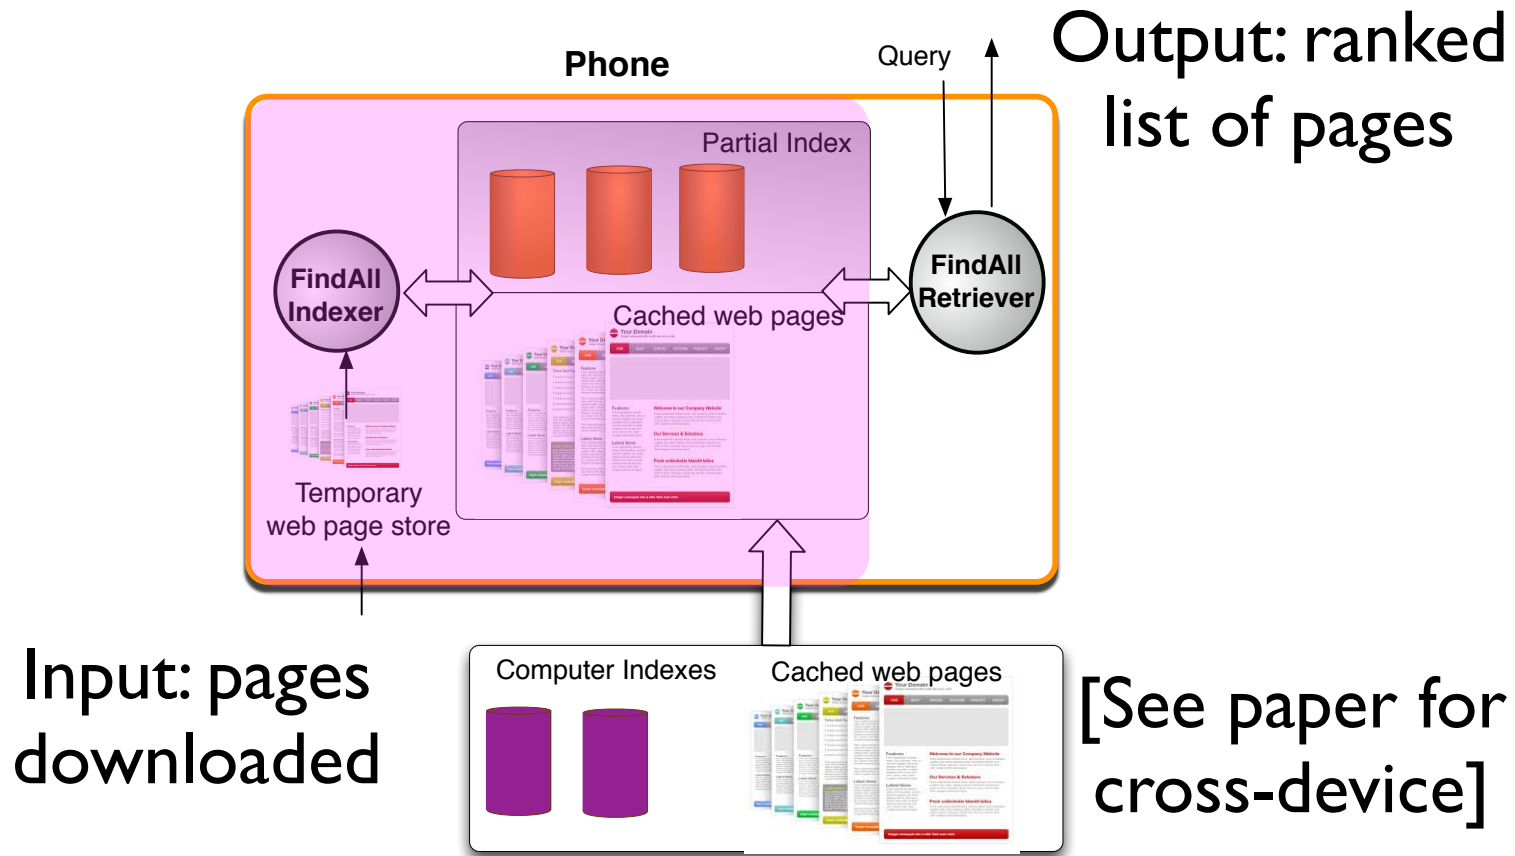
2. FindAll

- Design of search engine for mobile re-finding

3. Evaluation


- Results of the tradeoffs in practice

FindAll architecture



- Indexes as pages are browsed for availability
- Merges partial indexes to save energy

Key question: When to index?



High availability
High index energy

Block size
↔

Low availability
High retrieval energy

Tradeoff depends on user's behavior

Indexing strategy

- Index to maximize availability with energy consumption no higher than default search
- Index when:

$$\begin{array}{ccc} E[I] & & E[\neg I] \\ \text{Energy of indexing} & < & \text{Energy of default search} \end{array}$$

- Estimates need to be based on user behavior

Energy estimates [see paper]

- Energy of indexing, $E[I]$:
 - Cost of current block plus future penalty
- Energy of default search, $E[\neg I]$:
 - Sum of page download x re-find prob. of page
- Train a classifier for re-find prob. of page
 - Depends on several user features

Prototype on Android OS

- Adapt Galago search engine for phones
 - Implement partial indexing
- Implement online energy cost estimator
 - Train classifier when mobile is charging
 - Make an indexing decision every 5 mins
- Runs on a Motorola DroidX mobile
 - Has WiFi and Verizon 3G
 - Energy from Monsoon power monitor

Outline

1. User study

- Identifies key re-finding behaviors

2. FindAll

- Design of search engine for mobile re-finding

3. Evaluation


- Results of the tradeoffs in practice

Experiments

- Benefits and Costs
 - Latency, 3G data usage, availability
 - Energy, storage
- Alternate approaches
 - Cache keywords, query database
- Alternate indexing strategies
 - Cloud index, Always index, Fixed index


Results based on prototype and user trace

FindAll improves search latency



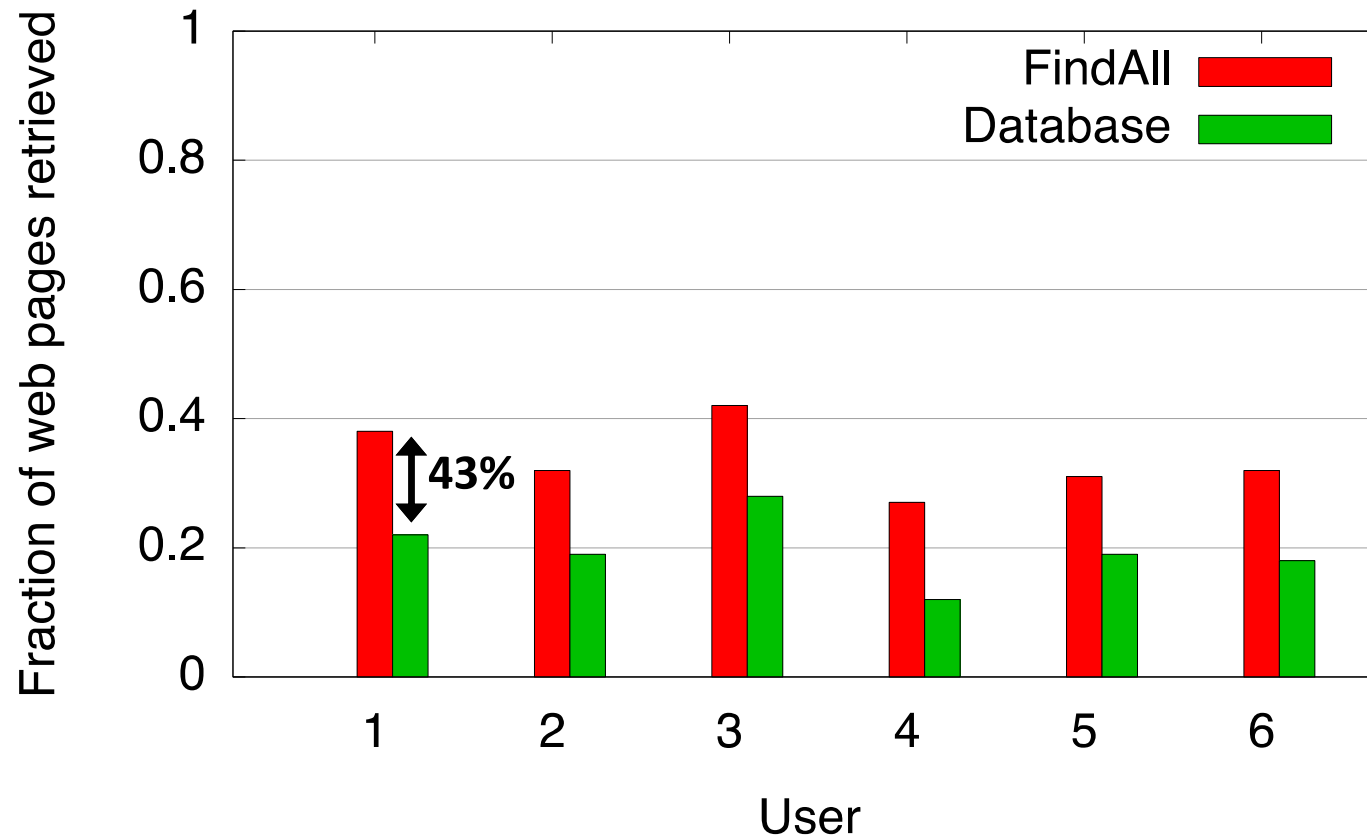
Average up to 2-fold better for high re-finders

FindAll reduces energy costs



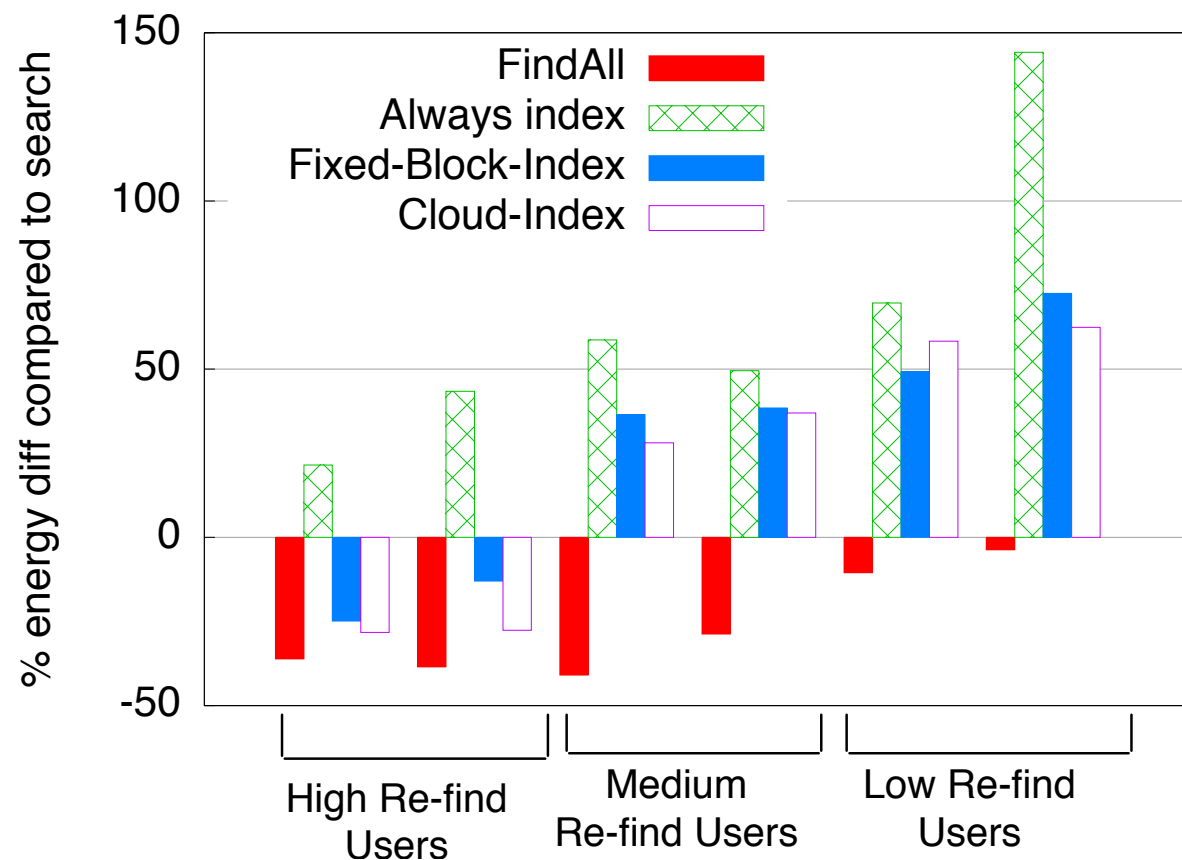
Energy up to 30% lower for high re-finders

FindAll improves availability when disconnected



(Under a random 50% connectivity model)

FindAll indexing strategy is important for energy benefits



Plus better availability than Cloud/Fixed

Conclusions

- FindAll makes a win-win tradeoff for search
 - Can decrease latency and increase availability, all with reduced energy and bandwidth
 - But need to adapt to user behavior
- Future directions
 - Integrating re-finding with mobile apps
 - Prediction across devices and users

Thank you. Questions?